

Performance V3 — Feature Spec

Product area: CRM Marketing

PM: M. Bourguignon

Fidelity: Near-production

Platform: Web (Arenamatrix app)

Date: 2026-06-10

△ **Agent notes:** The following were inferred during the interview and should be validated before build — the struggle magnitude; the cross-cutting guidelines (multi-channel low-complexity, uneven tracking fidelity, no-reload, partial-failure tolerance); the explicit non-goals; every story's happy path and edge states; the emotional driver behind the core job; and four derived mock-data entities (daily sales series, revenue/conversions time-series, channel composition aggregate, conversion rules config). All four jobs-to-be-done are currently tagged **assumed** (0 of 4 observed) — treat real marketing managers' first reactions as the validation gate.

1. Context

Performance V3 gives marketing managers (the core ICP) a rich page to analyze campaign performance across channels — Email, SMS, WhatsApp, Wallet — covering engagement stats (opens, clicks...) and, critically, ROI: the impact on sales after a campaign is sent. Today they have only the V2 page, which is too bare: it surfaces a single revenue-conversion KPI with little transparency on how it's computed and lacks the deeper KPIs and analysis managers need. Their workaround is exporting to spreadsheets or relying on ESP-native stats and flying blind on true cross-channel ROI. We're building it now because users need stronger ROI features to understand — and justify — the value proposition of the CRM.

Struggle magnitude: Weekly to several times/week · high — managers fly blind on ROI and can't defend CRM spend at budget reviews.

2. Goals

1. From any campaign, the manager sees its conversion into revenue (attributed revenue, €) — per campaign and overall.
2. The manager can configure how a conversion is calculated — choosing the trigger and the time range — within a bounded set of options (not free-form), so the page stays scalable. (Admin-only.)

3. The manager can view detailed engagement stats (opens, clicks, ...) per campaign and per channel to understand performance and adapt strategy.
4. The page provides a complete per-campaign detail view (in-page drawer).

Not in scope

- Per-recipient / individual-contact conversion list (this is an aggregated page).
- Actioning from the page beyond logging — no editing/duplicating/sending of existing campaigns (read & analysis only in V3).
- Fully free-form attribution config — only bounded trigger/range options.

3. User stories

Story 1

As a...	marketing manager
I want to...	see revenue + engagement across all my campaigns under a transparent conversion rule
So that...	I can judge my overall strategy
Acceptance criteria	Headline KPI strip: revenue (€ overall, + period delta), conversion %, open %, click % (degrade per channel availability); Engagement stats split by channel sub-tabs (Email/SMS/WhatsApp/Wallet), each showing only supported metrics + quality label + delta + trend; Conversions shown as one unified attributed-revenue total decomposed by channel (composition donut/stacked); Conversion rule (trigger + range) configurable within bounded options (admin-only; disabled for clients); Changing the rule/filters updates figures with no page reload (rule recompute may be async with a recomputing... state); Separate Tab 2 visual: campaigns marked on the general sales-trend curve
Empty state	No campaigns in this period — prompt to adjust filters or Log a campaign; KPIs render as €0 / —
Error state	Data fetch failed → retry; partial channels still render
Loading state	Skeleton KPI strip + charts

Story 2

As a...	marketing manager
I want to...	open one campaign's detail to see its stats and conversions
So that...	I understand why it performed
Acceptance criteria	Drawer shows dedicated KPIs: recipients, opens/clicks (or n/a per channel), conversion %, attributed revenue; Shows campaign info: subject, segment, send time, and the conversion rule applied; Close via X or click-outside; page behind is dimmed
Empty state	No conversions yet → show engagement stats, revenue €0 with a note
Error state	Detail load failed → retry
Loading state	Skeleton drawer

Story 3

As a...	marketing manager
I want to...	log a manual/offline campaign
So that...	it's analyzed alongside online ones
Acceptance criteria	Modal captures name, channel/type, date sent, reach/recipients, optional cost; Bounded fields only (no free-form attribution); On save it appears in the chronological banner list and rolls into aggregates; Engagement metrics render as n/a (offline) and are excluded from rate averages, manual revenue counts toward the total; Manual campaigns can later be edited/deleted (delete confirmed), online campaigns allow only limited field edits (description + a few)
Empty state	n/a (form)
Error state	Save failed → keep form data, show retry
Loading state	Saving spinner on submit

4. Screens

Screen	Linked stories	Notes
Performance (Tab 1)	1, 2, 3	Scrollable analytics page with sub-tabs and a side drawer: title row (Create campaign primary CTA · Options ▾ Import/Synchronise · Conversion rules dialog, admin-only) · controls (search + période, global across tabs) · headline KPI strip · engagement stats by channel sub-tabs · unified conversions with channel composition · paginated chronological campaign banner list, each with a : menu · detail drawer
Sales impact (Tab 2)	1	General sales-trend curve with campaign send markers (correlation, not attribution); hover a marker → campaign name/channel/revenue; compact campaign list for the selected période. Kept deliberately separate from attributed revenue.
Log campaign (modal)	3	Form: name, channel, date sent, reach, optional cost; bounded fields; Cancel / Save; on save appears in list + aggregates
Conversion rules (dialog, admin)	1	Per channel: trigger (Purchase/Visit...) + range (3/5/7 days) + small action set; disabled for clients, editable by admins; save triggers async recompute with a recomputing... state

Wireframe references

None specified (ASCII wireframes captured in the interview; Klaviyo and Brevo reporting screens used as visual inspiration).

Mock data

Entity	Key fields	Volume	Notes
Campaign	id, name, channel (Email/SMS/WhatsApp/Wallet/manual), date_sent, recipients, opens(count+rate), clicks(count+rate), conversions(count+rate), attributed_revenue(€), conversion_rule_applied, is_manual, editable_fields, subject, segment, send_time, cost(manual only)	20 (10 Email · 4 SMS · 2 WhatsApp · 1 Wallet · 3 Manual)	Edge: very long name; zero-revenue; high-reach ~27 700 + tiny ~37; negative open-rate delta; SMS/Wallet "not tracked" metrics; manual campaign with revenue but no engagement
Daily sales series	date, total_sales(€)	~90 days	Spike near a send + a flat/declining stretch so correlation reads honestly
Revenue/conversions time-series	bucket, revenue(€), conversions, split by channel	weekly buckets	One bucket with zero sends → gap in the bars
Channel composition aggregate	channel, revenue(€), share(%)	4 channels	Wallet share ~0% → near-empty slice
Conversion rules	channel, trigger (Purchase/Visit...), range (3/5/7 days), action_set	4 channel rows	One channel on a non-purchase trigger (Visit)

5. Constraints

Stack / framework	Arenamatrix design system
Accessibility	WCAG AA
Data & GDPR	None specified
Other	Good and scalable performance; conversion-rule recompute may be asynchronous

Open questions

Question	Owner	Status
Are the four jobs-to-be-done real? All tagged "assumed" (0/4 observed) — validate with marketing managers, starting with the core job (transparent attributed revenue)	M. Bourguignon	Open
How will we know Job 3 (log offline campaign) is actually used? Success signal is soft ("completeness of channel mix")	M. Bourguignon	Open
Kill criteria on record: if attribution is too unreliable/sparse to show trustworthy ROI, or we can't surface enough metrics for managers to understand their activity — don't ship over V2	M. Bourguignon	Open

6. Pre-send checklist

- ✓ Context is one clear paragraph — who, what, why
- ✓ At least 1 goal uses outcome language
- ✓ At least 1 non-goal is stated explicitly
- ✓ Every story has acceptance criteria or an explicit note that they are missing
- ✓ All 3 edge states are specified for every story
- ✓ Every screen is listed with a description or wireframe reference
- ✓ At least 1 mock data entity is defined
- ✓ All open questions have an owner